

Big data science Day 3



F. Legger - INFN Torino


<https://github.com/Course-bigDataAndML/MLCourse-2324>

Yesterday

- Big data, analytics
- Distributed computing
- ML: Feature engineering

Today

- **Machine learning**
 - Architectures
 - Train model and evaluate



Remember:




Machine learning involves
two mathematical entities



Model: a mathematical model
describes the relationship between
different aspects of the data



Features: a feature is a
representation of raw data



Model



Feature engineering

Raw Data

```
0 : {  
  house_info : {  
    num_rooms: 6  
    num_bedrooms: 3  
    street_name: "Shorebird Way"  
    num_basement_rooms: -1  
    ...  
  }  
}
```

Raw data doesn't come to us as feature vectors.

Feature Engineering

Feature Vector


```
[  
  6.0,  
  1.0,  
  0.0,  
  0.0,  
  0.0,  
  9.321,  
  -2.20,  
  1.01,  
  0.0,  
  ...  
]
```

Process of creating features from raw data is **feature engineering**.

Example: supervised classification

Ingredients

- **Inputs:** \mathbf{X} , is a matrix of size \mathbf{n} (number of samples) \times \mathbf{m} (number of features)
- **Features:** \mathbf{X} , transformed inputs, matrix $\mathbf{n} \times \mathbf{m}$
- **Labels:** \mathbf{y} , vector size \mathbf{n}



Recipe: supervised classification

- For each input vector X_i predict z_i , $i=1\dots n$




- **Predictions:**
 - $z_i = \phi(W^T X_i)$ yields (0,1)
 - **Weights:** W (matrix) contains the model parameters
 - **Activation function:** ϕ (step function, sigmoid)
- **Cost function == loss function == prediction error**, function of the model parameters W
- **Aim: find weights W that minimize cost function**

Activation function

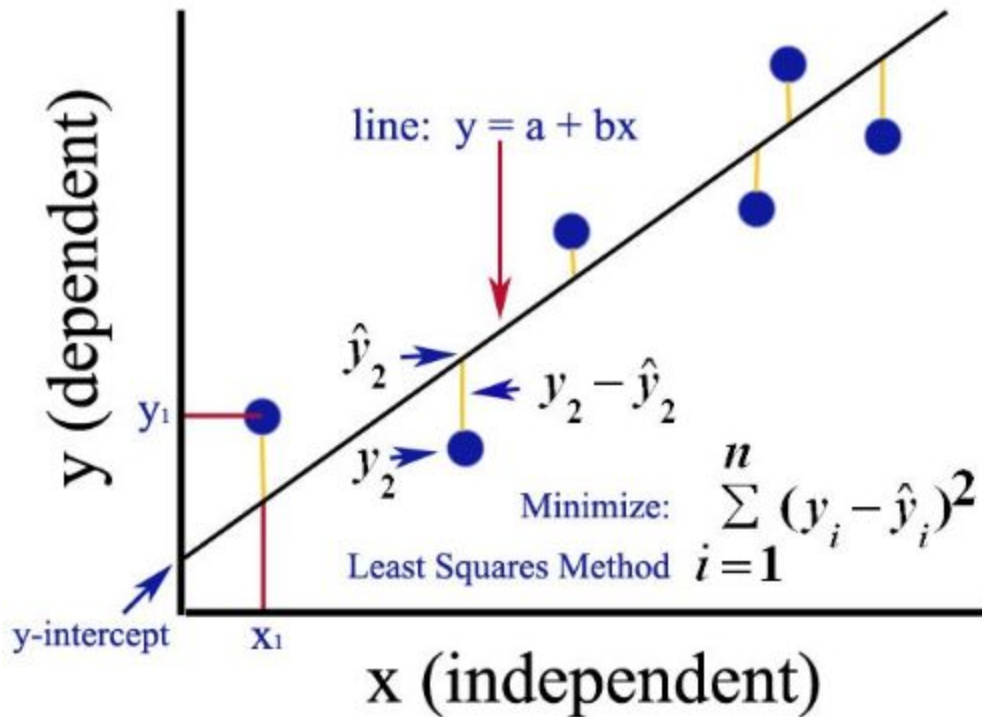
- Turns unbounded output into a known range/shape
- For example, **sigmoid** function only outputs numbers in the range (0, 1)
 - big negative numbers become ~0
 - big positive numbers become ~1.

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}.$$



Another example, linear regression

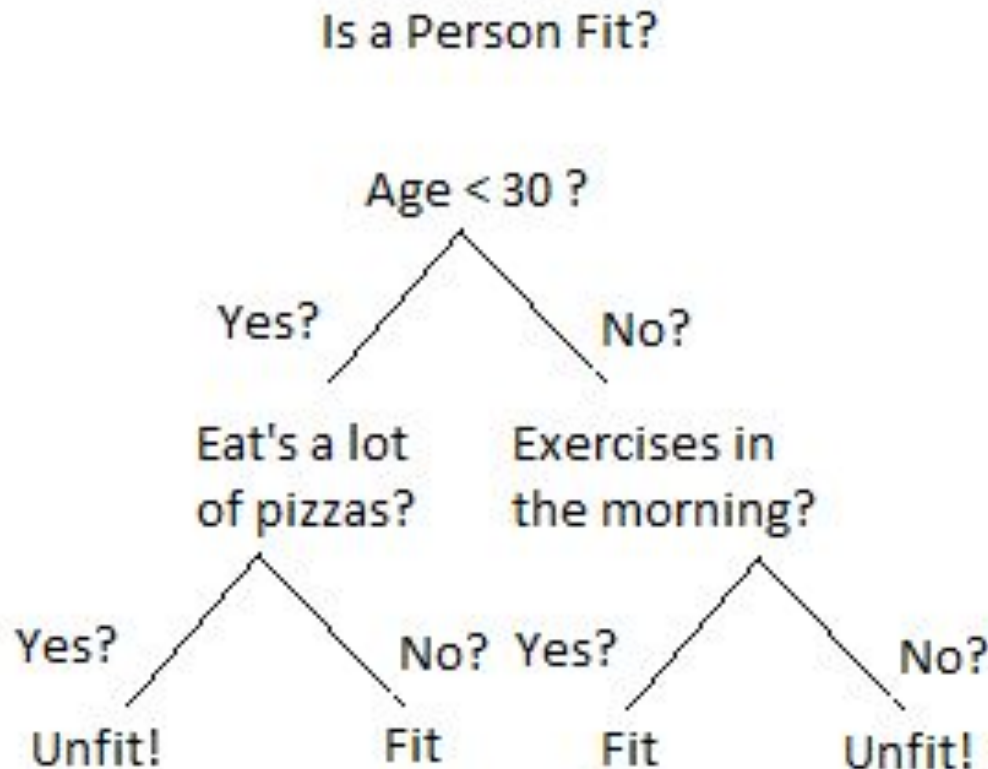
- Inputs (features): x_i
- Labels: y_i
- Model: $y = a + bx$
- Weight+bias (parameters to be found): a, b
- Cost function: **Mean Square Error (MSE)**
- No **activation function**: problem is linear



$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

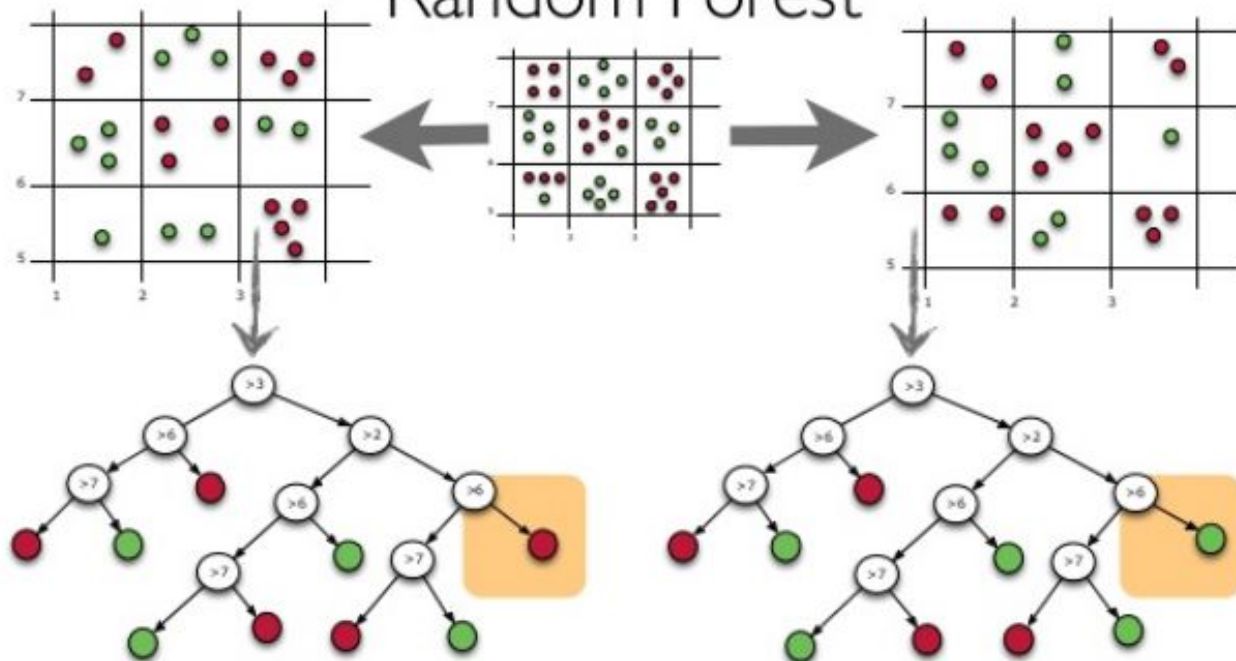
Supervised learning

Decision trees: supervised classification



Typically used in combinations (Random forest, Gradient Tree Boosting)

Random Forest



- Each tree sees part of the training sets and captures part of the information it contains







Ensembles

- **Bagging**
 - building multiple models (typically of the *same* type) from different subsamples of the training dataset
- **Boosting**
 - building multiple models (typically of the *same* type) each of which learns to fix the predictions errors of a prior model in the chain
- **Stacking**
 - building multiple models (typically of *different* types) and a supervisor model that learns how to best combine the predictions of the primary model
- **Weighting|Blending**
 - combine multiple models into single prediction using different weight functions



Neural networks: supervised classification

- Basic unit: **Neuron**. A neuron takes inputs, does some math with them, and produces **one** output

More on NN tomorrow!



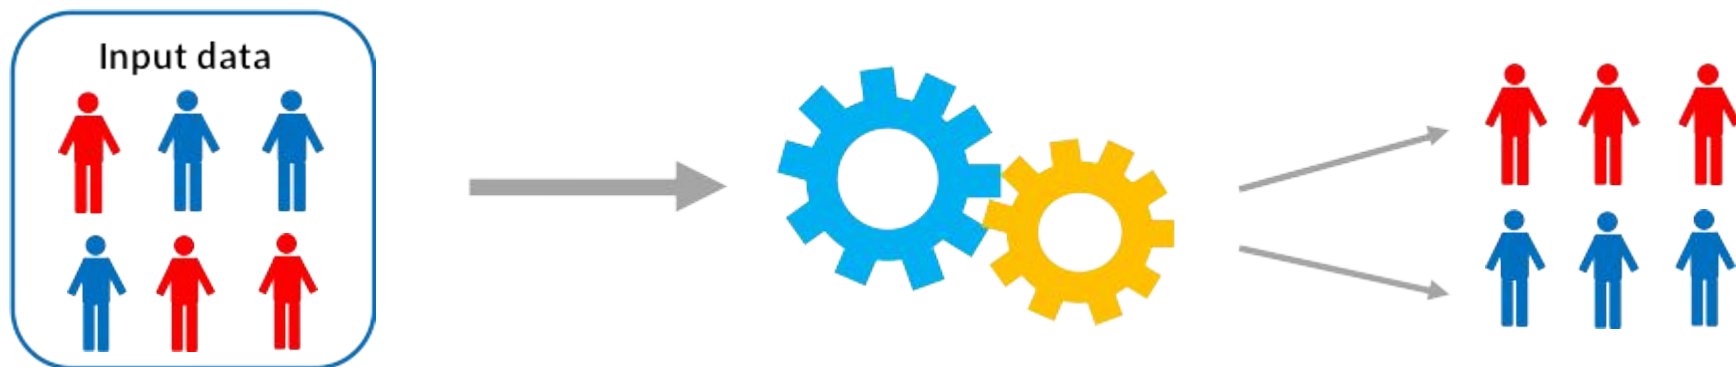
Support vector machines (SVG), supervised classification




Unsupervised learning

Challenges


- No label (ground truth) in input dataset
- The system must have the ability to recognize patterns in the data without explicitly being told what patterns to identify




Principal Component Analysis (PCA), unsupervised




K-means clustering, unsupervised




1. k initial "means" (in this case $k=3$) are randomly generated within the data domain (shown in color).



2. k clusters are created by associating every observation with the nearest mean. The partitions here represent the **Voronoi diagram** generated by the means.



3. The **centroid** of each of the k clusters becomes the new mean.






4. Steps 2 and 3 are repeated until convergence has been reached.

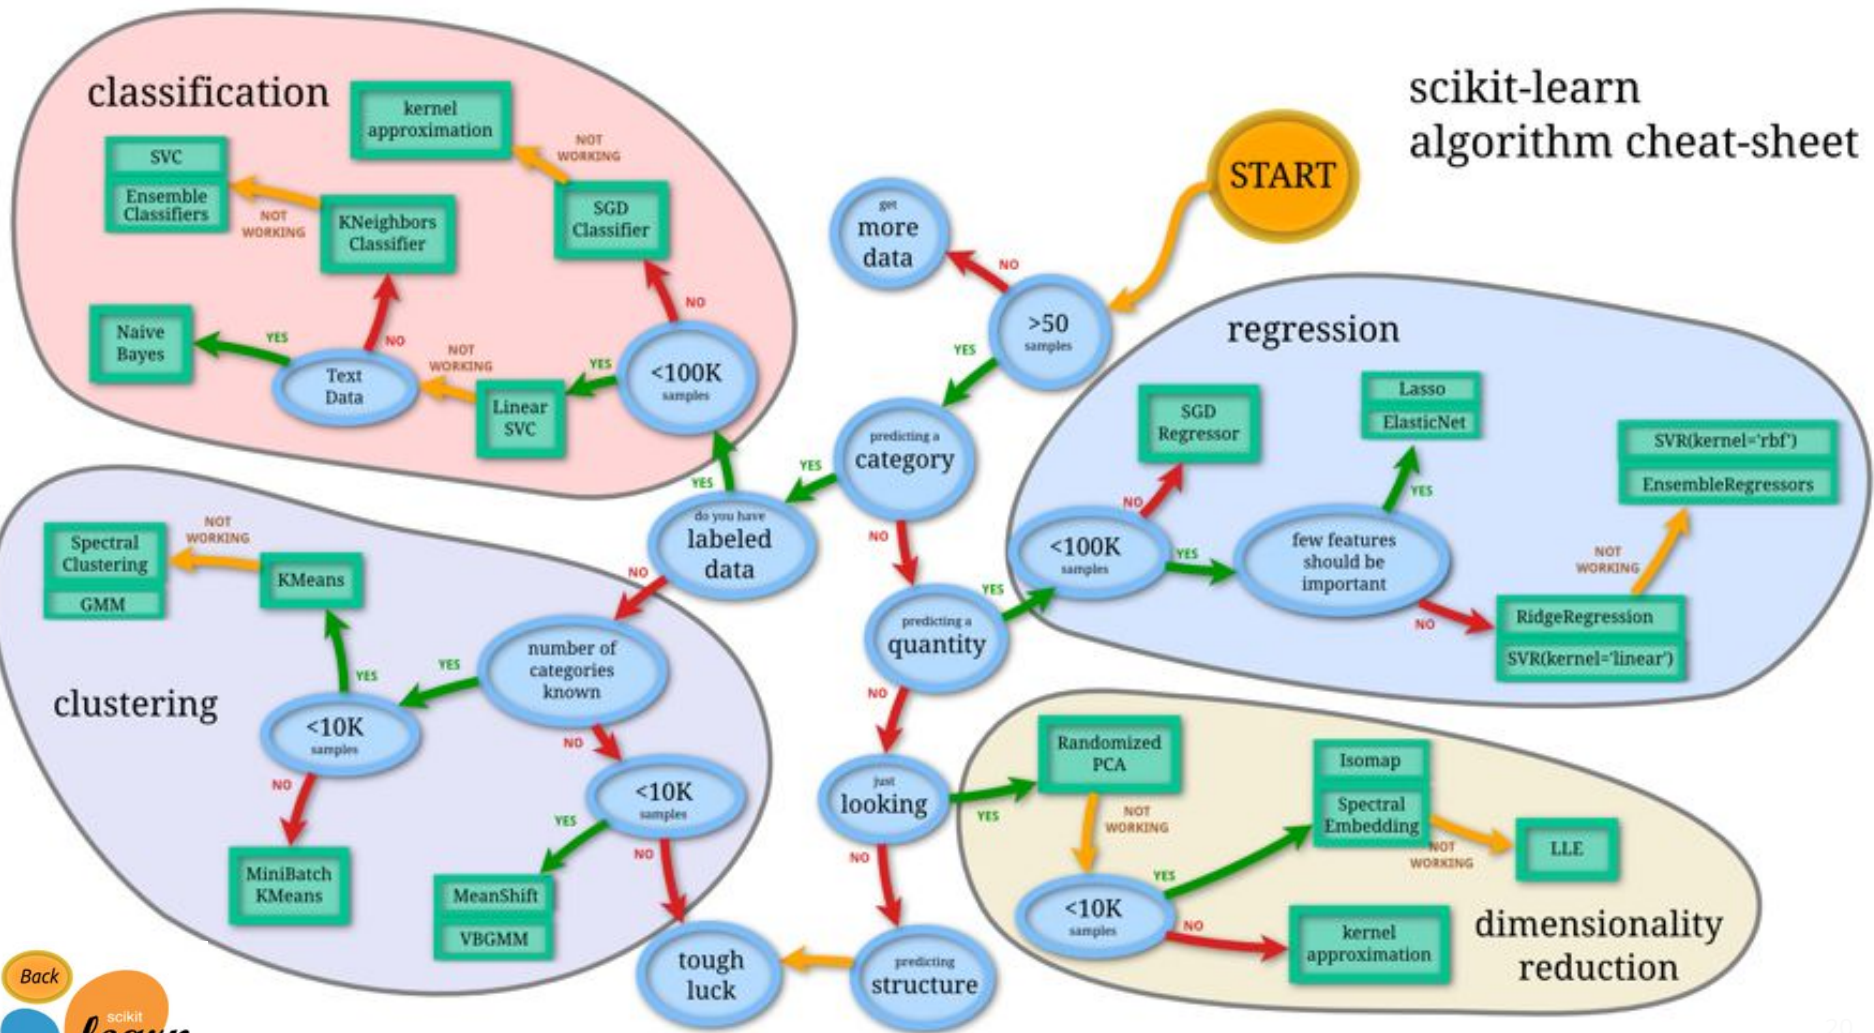
You must define **k**, the number of clusters, and which **distance** to use!







K-means: Distances

Taxicab or **Manhattan** distance:
sum of the projections along all
axis



scikit-learn algorithm cheat-sheet

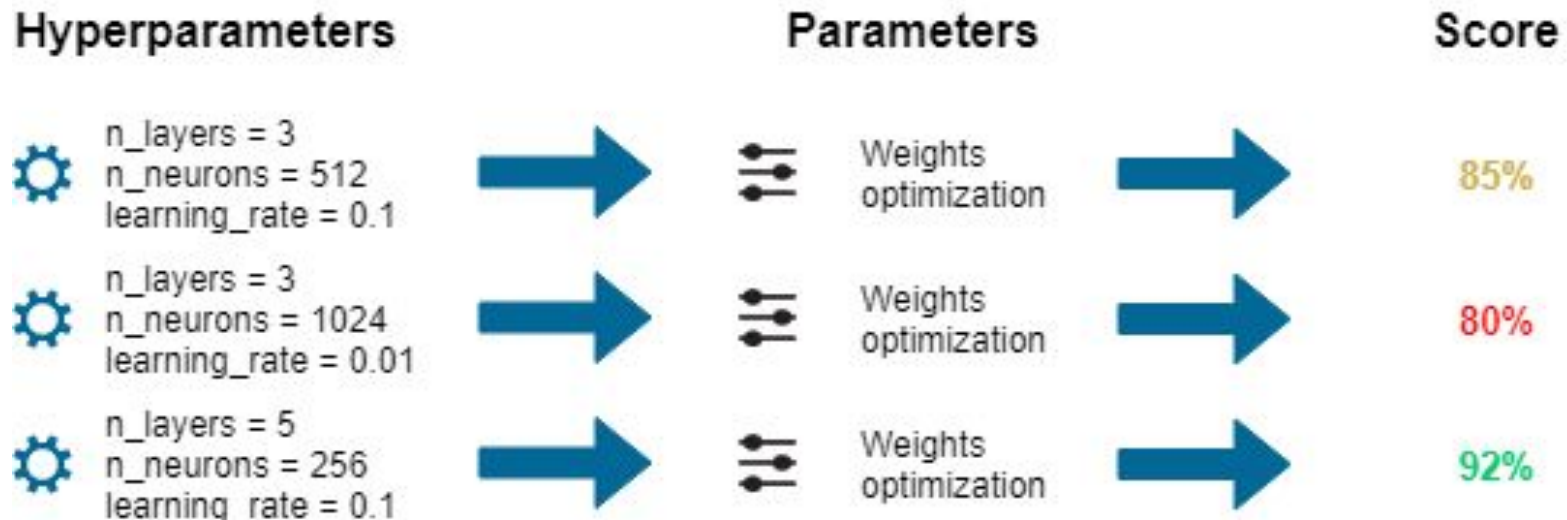


	TYPE	NAME	DESCRIPTION	ADVANTAGES	DISADVANTAGES
Linear		Linear regression	The “best fit” line through all data points. Predictions are numerical.	Easy to understand – you clearly see what the biggest drivers of the model are.	<ul style="list-style-type: none"> ✗ Sometimes too simple to capture complex relationships between variables. ✗ Tendency for the model to “overfit”.
		Logistic regression	The adaptation of linear regression to problems of classification (e.g., yes/no questions, groups, etc.)	Also easy to understand.	<ul style="list-style-type: none"> ✗ Sometimes too simple to capture complex relationships between variables. ✗ Tendency for the model to “overfit”.
Tree-based		Decision tree	A graph that uses a branching method to match all possible outcomes of a decision.	Easy to understand and implement.	<ul style="list-style-type: none"> ✗ Not often used on its own for prediction because it's also often too simple and not powerful enough for complex data.
		Random Forest	Takes the average of many decision trees, each of which is made with a sample of the data. Each tree is weaker than a full decision tree, but by combining them we get better overall performance .	A sort of “wisdom of the crowd”. Tends to result in very high quality models. Fast to train.	<ul style="list-style-type: none"> ✗ Can be slow to output predictions relative to other algorithms. ✗ Not easy to understand predictions.
		Gradient Boosting	Uses even weaker decision trees, that are increasingly focused on “hard” examples .	High-performing.	<ul style="list-style-type: none"> ✗ A small change in the feature set or training set can create radical changes in the model. ✗ Not easy to understand predictions.
Neural networks		Neural networks	Mimics the behavior of the brain. Neural networks are interconnected neurons that pass messages to each other. Deep learning uses several layers of neural networks put one after the other.	Can handle extremely complex tasks - no other algorithm comes close in image recognition.	<ul style="list-style-type: none"> ✗ Very, very slow to train, because they have so many layers. Require a lot of power. ✗ Almost impossible to understand predictions.


*All models are wrong, but
some are useful (George Box)*

Hyperparameters vs parameters


- Model **parameters** are learned during **training** when we optimize a loss function -> *weights*
- **Hyperparameters** are not model parameters and they **cannot** be directly trained from the data -> *model architecture*



Hyperparameter tuning



(a) Standard Grid Search



(b) Random Search

Modeling Algorithm

Tune

hyperparameters (tuning options)

- Polynomial order, penalty parameter, ...
- Network configuration, solver options, ...
- Max tree depth, splitting criterion, ...

Model




Train

model parameters

- Regression coefficients
- Neural net weights
- Tree splitting rules

...

Overfitting / underfitting



Underfitting

Model doesn't have enough (hyper-)parameters to describe data




Overfitting

Model has too many (hyper-)parameters

Classification metrics

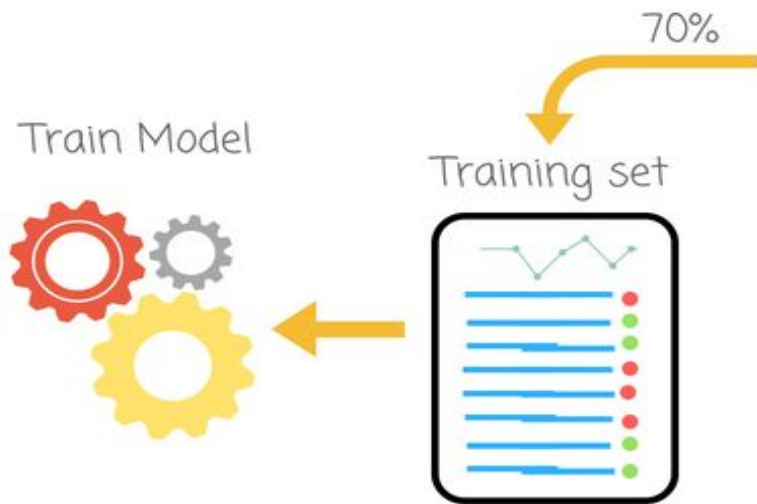
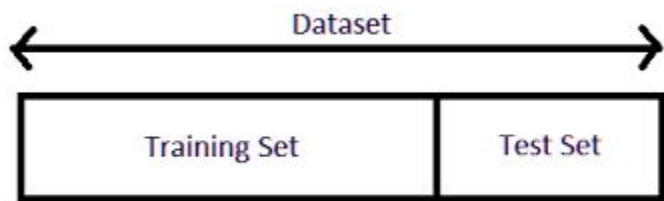
- **ROC**: Receiver Operating Characteristics
- **AUC**: Area under the curve
- **TPR**: True positive rate
- **FPR**: False positive rate
- **TNR/FNR**: True/False negative rate




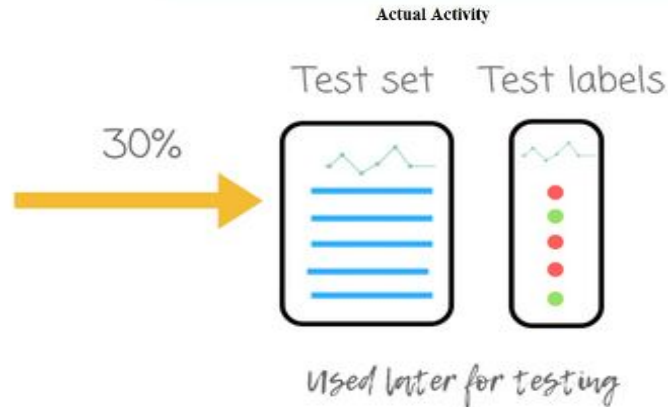
		True class			
		p	n		
Hypothesized class	Y	True Positives	False Positives	$fp\ rate = \frac{FP}{N}$	$tp\ rate = \frac{TP}{P}$
	N	False Negatives	True Negatives	$precision = \frac{TP}{TP+FP}$	$recall = \frac{TP}{P}$
Column totals:		P	N	$accuracy = \frac{TP+TN}{P+N}$	
				$F\text{-measure} = \frac{2}{1/precision + 1/recall}$	

Confusion matrix

Training and test set


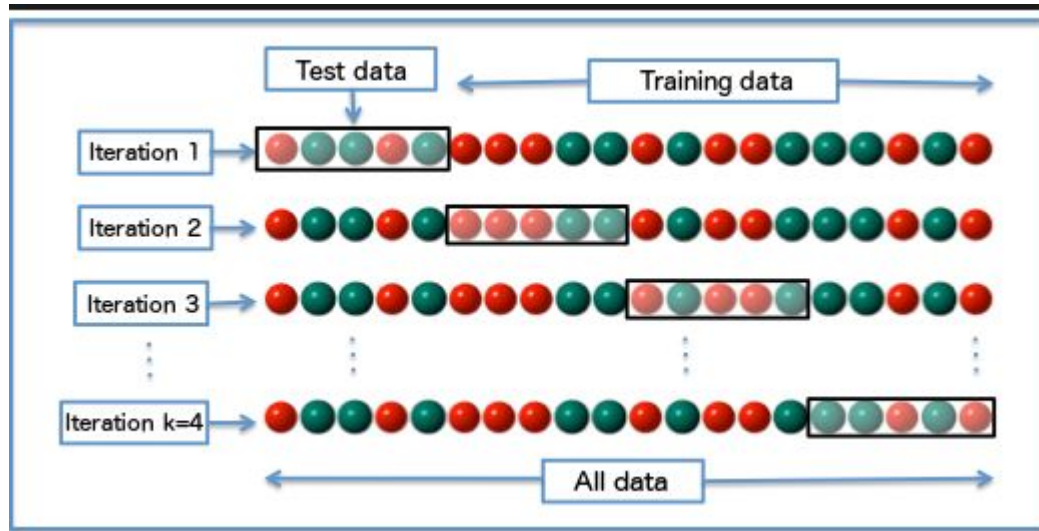


Entire Dataset

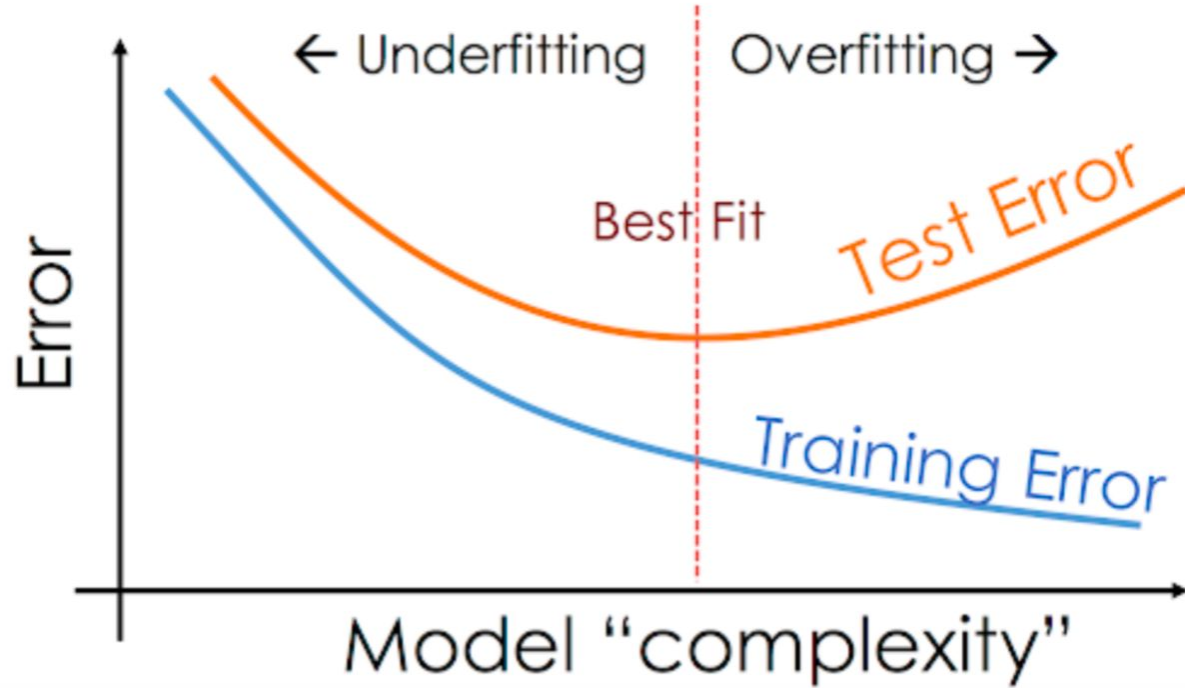


Cross-validation: is your model robust?

- Train/test split
- K-folds cross validation



Loss, or could
be any other
metrics of
interest



Or training epochs (number of iterations)

Uncertainties

- **Covariate shift/domain shift:** change in the distribution of the input variables between the training and the test datasets
- **How to treat:**
 - **Propagate all errors** in your input variables to the output model
 - Use **cross-validation** and get an estimate of the variance of the model performances when changing hyperparameters
 - Compare model performances when **dropping/adding** features
 - If possible, compare your model results with another architecture
 - Ideally, measure error with independent model-data comparison (such as tag&probe method) just like any other classifier
- [Reference](#)